



Two-account layout — A improves & pseudonymizes (all versions); B holds clean, deduplicated *slice snapshots* (no continuous history)

AWS Account A — Improvement (bitemporal) zone (all versions; PII & re-id map live here)



account
boundary

AWS Account B — Pseudonymized zone (clean snapshots · no PII · no continuous history)



What Account A does — all of it. A is the improvement zone: the bitemporal variant already holds *pseudonymized*, quality-corrected, imputed, defect-corrected data. Pseudonymization is just one of A's improvements, not a boundary step. A keeps *all* versions on both axes, including ones that are *not yet clean*.

The A → B transform is a projection, not an improvement. It (1) *strips errant / not-yet-clean versions* — the reason A's full variant can't simply land in B is that some versions haven't passed correction yet — and (2) *collapses the SCD2 version explosion* — the explosion that combining tenants creates in the first place — by materializing a chosen temporal slice with survivorship. So B has no continuous history — no incremental changelog to walk — only deduplicated *snapshots*, though any slice can be generated and snapshotted on demand. In the normal path, only pseudonymized, clean data crosses and the re-id map stays in A — but, by *exception* (a policy decision), an authorized B consumer may be granted audited re-identification: results served from A, or — policy permitting — direct access to the re-id map itself.